

In [1]:

```
import pandas as pd
import os
os.getcwd()
```

Out[1]:

```
'C:\\Users\\Mahima Sharu'
```

In [2]:

```
os.listdir(os.getcwd())
```

Out[2]:

```
[ '.anaconda',
  '.android',
  '.AndroidStudio3.2',
  '.bash_history',
  '.conda',
  '.condarc',
  '.config',
  '.gitconfig',
  '.idlerc',
  '.ipynb_checkpoints',
  '.ipython',
  '.jupyter',
  '.matplotlib',
  '.node_repl_history',
  '.packettracer',
  '.PyCharmCE2019.1',
  '.VirtualBox',
  '.vscode',
  '3 - 1st and 2nd.ipynb',
  '3 - 3rd.ipynb',
  '3 - 4th.ipynb',
  '3D Objects',
  '4th - 1st.ipynb',
  'Amazon Reviews_ Unlocked Mobile Phones _ Kaggle.csv',
  'Amazon Reviews_ Unlocked Mobile Phones _ Kaggle_files',
  'AppData',
  'Application Data',
  'assignment 1.1.ipynb',
  'assignment 1.3.txt',
  'assignment1.4.ipynb',
  'Cisco Packet Tracer 7.2',
  'Contacts',
  'Cookies',
  'debug.log',
  'Desktop',
  'Documents',
  'Downloads',
  'electric motor temperature.ipynb',
  'Favorites',
  'file.txt',
  'Git-workshop',
  'Hotel booking demand _ Kaggle.csv',
  'Hotel booking demand _ Kaggle_files',
  'HP',
  'index.html',
  'input.txt',
  'IntelGraphicsProfiles',
  'iris.csv',
  'KCLT.csv',
  'Links',
  'Local Settings',
  'MicrosoftEdgeBackups',
  'Music',
  'My Documents',
  'NetHood',
  'NTUSER.DAT',
  'ntuser.dat.LOG1' ]
```

```

'ntuser.dat.LOG2',
'NTUSER.DAT{43757b42-e0ed-11e9-986f-e08e1ad91340}.TM.blf',
'NTUSER.DAT{43757b42-e0ed-11e9-986f-e08e1ad91340}.TMContainer000000000000000001.regtrans-ms',
'NTUSER.DAT{43757b42-e0ed-11e9-986f-e08e1ad91340}.TMContainer000000000000000002.regtrans-ms',
'ntuser.ini',
'OneDrive',
'Oracle',
'Pictures',
'pmsm_temperature_data.csv',
'PrintHood',
'PycharmProjects',
'Recent',
'Saved Games',
'ScStore',
'Searches',
'SendTo',
'Sentiment Analysis on Movie Reviews _ Kaggle.csv',
'Sentiment Analysis on Movie Reviews _ Kaggletest.csv',
'Sentiment Analysis on Movie Reviews _ Kaggletest_files',
'Sentiment Analysis on Movie Reviews _ Kaggletrain.csv',
'Sentiment Analysis on Movie Reviews _ Kaggletrain_files',
'Sentiment Analysis on Movie Reviews _ Kaggle_files',
'Start Menu',
'Templates',
'Untitled.ipynb',
'Untitled1.ipynb',
'Untitled10.ipynb',
'Untitled11.ipynb',
'Untitled12.ipynb',
'Untitled2.ipynb',
'Untitled3.ipynb',
'Untitled4.ipynb',
'Untitled5.ipynb',
'Untitled6.ipynb',
'Untitled7.ipynb',
'Untitled8.ipynb',
'Untitled9.ipynb',
'Videos',
'VirtualBox VMs',
'wekafiles',
'_split.csv']

```

In [9]:

```

#Load File
df=pd.read_csv('C:\\Users\\Mahima Sharu\\pmsm_temperature_data.csv')

```

Out[9]:

	ambient	coolant	u_d	u_q	motor_speed	torque	i_d	i_q	pm	stator_yoke	stator_tooth	stator
0	0.752143	1.118446	0.327935	1.297858	-1.222428	0.250182	1.029572	0.245860	2.522071	-1.831422	-2.066143	
1	0.771263	1.117021	0.329665	1.297686	-1.222429	0.249133	1.029509	0.245832	2.522418	-1.830969	-2.064859	
2	0.782892	1.116681	0.332771	1.301822	-1.222428	0.249431	1.029448	0.245818	2.522673	-1.830400	-2.064073	
3	0.780935	1.116764	0.333700	1.301852	-1.222430	0.248636	1.032845	0.246955	2.521639	-1.830333	-2.063137	
4	0.774043	1.116775	0.335206	1.303118	-1.222429	0.248701	1.031807	0.246610	2.521900	-1.830498	-2.062795	
...
998065	0.047497	0.341638	0.331475	1.246114	-1.222428	0.255640	1.029142	0.245722	0.429853	1.018568	0.836084	
998066	0.048839	0.320022	0.331701	1.250655	-1.222437	0.255640	1.029148	0.245736	0.429751	1.013417	0.834438	
998067	0.042350	0.307415	0.330946	1.246852	-1.222430	0.255640	1.029191	0.245701	0.429439	1.002906	0.833936	
998068	0.039433	0.302082	0.330987	1.249505	-1.222432	0.255640	1.029147	0.245727	0.429558	0.999157	0.830504	
998069	...	0.312666	0.330830	...	-1.222431	...	1.029141	...	0.429166	0.987163	0.828046	

```
0.043803 ambient coolant u_d 1.246591 u_q motor_speed 0.255640 torque i_d 0.245722 i_q pm stator_yoke stator_tooth stator
```

998070 rows × 13 columns

In [4]:

```
from sklearn.preprocessing import Imputer
from sklearn.model_selection import KFold #Provides train/test indices to split data in
train/test sets
from sklearn import linear_model
from sklearn.metrics import make_scorer

from sklearn import svm #Support Vector Machine are a set of supervised learning methods
# linear algebra
import numpy as np
# data processing, CSV file I/O (e.g. pd.read_csv)
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns #Visualizing dataset structure that can be used to make visualizations
with multiple plots
from sklearn.linear_model import LinearRegression

from sklearn import neighbors
from math import sqrt
```

In [10]:

```
*****Basic Linear
Regression*****

import statsmodels.api as sm #provides classes and functions for the estimation of many
different statistical models

#Defining dependent and independent variable
X = df['i_d']
X=sm.add_constant(X)

y = df['motor_speed']

lm=sm.OLS(y,X) #Leastsquare Minimization using Ordinary Least Square value along the
x and y axes
model=lm.fit() #Data Fitting

model.summary()
```

C:\ProgramData\Anaconda3\lib\site-packages\numpy\core\fromnumeric.py:2389: FutureWarning: Method .ptp is deprecated and will be removed in a future version. Use numpy.ptp instead.
return ptp(axis=axis, out=out, **kwargs)

Out[10]:

OLS Regression Results

Dep. Variable:	motor_speed	R-squared:	0.523
Model:	OLS	Adj. R-squared:	0.523
Method:	Least Squares	F-statistic:	1.093e+06
Date:	Fri, 03 Apr 2020	Prob (F-statistic):	0.00
Time:	12:17:00	Log-Likelihood:	-1.0484e+06
No. Observations:	998070	AIC:	2.097e+06
Df Residuals:	998068	BIC:	2.097e+06
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	-0.0020	0.001	-2.826	0.005	-0.003	-0.001
i_d	-0.7245	0.001	-1045.267	0.000	-0.726	-0.723

Omnibus:	7561.278	Durbin-Watson:	0.003
Prob(Omnibus):	0.000	Jarque-Bera (JB):	5931.452
Skew:	-0.109	Prob(JB):	0.00
Kurtosis:	2.691	Cond. No.	1.01

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [11]:

```
model.params
```

Out[11]:

```
const    -0.001957
i_d      -0.724531
dtype: float64
```

In [12]:

```
print("f_pvalue:", "%.4f" % model.f_pvalue)
```

```
f_pvalue: 0.0000
```

In [13]:

```
#mean square value
model.mse_model
```

Out[13]:

```
522878.40071765514
```

In [14]:

```
model.rsquared
```

Out[14]:

```
0.5226043734324983
```

In [15]:

```
model.rsquared_adj
```

Out[15]:

```
0.522603895112758
```

In [16]:

```
#Predicted values
model.fittedvalues[0:5]
```

Out[16]:

```
0    -0.747914
1    -0.747869
2    -0.747824
3    -0.750285
4    -0.749534
dtype: float64
```

In [17]:

```
#Real values
y[0:5]
```

Out[17]:

```
0   -1.222428
1   -1.222429
2   -1.222428
3   -1.222430
4   -1.222429
Name: motor_speed, dtype: float64
```


In [18]:

```
#Model equation
print("Motor speed = " +
      str("%.3f" % model.params[0]) + ' + i_d' + "*" +
      str("%.3f" % model.params[1]))
```

Motor speed = -0.002 + i_d*-0.725

In [19]:

```
#Regression Model Visualization
g=sns.regplot(df['i_d'], df['motor_speed'],
              ci=None, scatter_kws={'color': 'r', 's':9})
g.set_title('Model equation: motor_speed = -0.002 + i_d * -0.725')
g.set_ylabel('Motor_speed')
g.set_xlabel('i_d');
```



In [20]:

```
from sklearn.metrics import r2_score, mean_squared_error
mse=mean_squared_error(y, model.fittedvalues)
rmse=np.sqrt(mse)
rmse
```

Out[20]:

0.6917872418443057

In [21]:

```
k_t=pd.DataFrame({'Real_values':y[0:50],
                  'Predicted_values':model.fittedvalues[0:50]})
k_t['error']=k_t['Real_values']-k_t['Predicted_values']
k_t.head()
```

Out[21]:

Real_values	Predicted_values	error
-------------	------------------	-------

0	-1.222428 Real_values	-0.747914 Predicted_values	-0.474514 error
1	-1.222429	-0.747869	-0.474561
2	-1.222428	-0.747824	-0.474604
3	-1.222430	-0.750285	-0.472145
4	-1.222429	-0.749534	-0.472895

In [22]:


```
#Easiest way to learn residual
model.resid[0:10]
```

Out[22]:

```
0    -0.474514
1    -0.474561
2    -0.474604
3    -0.472145
4    -0.472895
5    -0.473457
6    -0.473848
7    -0.474130
8    -0.474315
9    -0.474471
dtype: float64
```

In [23]:

```
plt.plot(model.resid);
```



In [24]:

```
*****Multiple Linear
Regression*****
X=df.drop("motor_speed", axis=1)
y=df["motor_speed"]
```

In [25]:

```
from sklearn.model_selection import train_test_split,cross_val_score,cross_val_predict
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2, random_state=42)
training=df.copy()
```

In [26]:

```
lm=sm.OLS(y_train, X_train)
model=lm.fit()
#All coefficients are significant for the model by looking at the p-value. ( P>|t| )
model.summary()
```

Out[26]:

OLS Regression Results

Dep. Variable:	motor_speed	R-squared (uncentered):	0.928
Model:	OLS	Adj. R-squared (uncentered):	0.928
Method:	Least Squares	F-statistic:	8.582e+05
Date:	Fri, 03 Apr 2020	Prob (F-statistic):	0.00
Time:	12:35:22	Log-Likelihood:	-83308.
No. Observations:	798456	AIC:	1.666e+05
Df Residuals:	798444	BIC:	1.668e+05
Df Model:	12		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
ambient	-0.0503	0.000	-131.429	0.000	-0.051	-0.050
coolant	0.4091	0.002	218.100	0.000	0.405	0.413
u_d	-0.1657	0.001	-254.611	0.000	-0.167	-0.164
u_q	0.5394	0.000	1469.606	0.000	0.539	0.540
torque	-0.3411	0.005	-70.287	0.000	-0.351	-0.332
i_d	-0.6580	0.001	-1268.145	0.000	-0.659	-0.657
i_q	0.1352	0.005	29.630	0.000	0.126	0.144
pm	0.1061	0.001	170.646	0.000	0.105	0.107
stator_yoke	-1.6278	0.006	-282.304	0.000	-1.639	-1.617
stator_tooth	2.3219	0.008	304.592	0.000	2.307	2.337
stator_winding	-1.1714	0.004	-310.271	0.000	-1.179	-1.164
profile_id	-1.117e-05	5.6e-06	-1.995	0.046	-2.22e-05	-1.96e-07

Omnibus:	43472.517	Durbin-Watson:	2.000
Prob(Omnibus):	0.000	Jarque-Bera (JB):	169112.669
Skew:	-0.111	Prob(JB):	0.00
Kurtosis:	5.244	Cond. No.	1.82e+03

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.82e+03. This might indicate that there are strong multicollinearity or other numerical problems.

In [27]:

```
#Root Mean Squared Error for Train
rmse1=np.sqrt(mean_squared_error(y_train,model.predict(X_train)))
rmse1
```

Out[27]:

0.2685809987166384

In [28]:

```
#Root Mean Squared Error for Test
rmse2=np.sqrt(mean_squared_error(y_test,model.predict(X_test)))
rmse2
```

Out[28]:

0.26823759191828583

In [29]:

```
#Model Tuning for Multiple Linear Regression
model = LinearRegression().fit(X_train,y_train)
cross_val_score1=cross_val_score(model, X_train, y_train, cv=10, scoring='r2').mean() #verified
score value for train model
print('Verified R2 value for Training model: ' + str(cross_val_score1))

cross_val_score2=cross_val_score(model, X_test, y_test, cv=10, scoring='r2').mean() #verified score
value for test model
print('Verified R2 value for Testing Model: ' + str(cross_val_score2))
```

Verified R2 value for Training model: 0.9280570425519296

Verified R2 value for Testing Model: 0.9281973950061891

In [30]:

```
#For Root Mean square value
RMSE1=np.sqrt(-cross_val_score(model, X_train, y_train, cv=10,
                                scoring='neg_mean_squared_error')).mean() #verified RMSE score value
for train model
print('Verified RMSE value for Training model: ' + str(RMSE1))

RMSE2=np.sqrt(-cross_val_score(model, X_test, y_test, cv=10,
                                scoring='neg_mean_squared_error')).mean() #verified RMSE score value
for test model
print('Verified RMSE value for Testing Model: ' + str(RMSE2))
```


Verified RMSE value for Training model: 0.2685584755640468

Verified RMSE value for Testing Model: 0.26822947422677207

In [31]:

```
#Visualizing for Multiple Linear Regression y values

import seaborn as sns
ax1 = sns.distplot(y_train, hist=False, color="r", label="Actual Value")
sns.distplot(y_test, hist=False, color="b", label="Fitted Values" , ax=ax1);
```



In [33]:

```
*****Principal Component
Regression*****
from sklearn.decomposition import PCA
from sklearn.preprocessing import scale


pca=PCA()
X_reduced_train=pca.fit_transform(scale(X_train))
```

In [34]:

```
explained_variance_ratio=np.cumsum(np.round(pca.explained_variance_ratio_ , decimals=4)* 100)[0:20]
```


In [35]:

```
plt.bar(x=range(1, len(explained_variance_ratio)+1), height=explained_variance_ratio)
plt.ylabel('percentage of explained variance')
plt.xlabel('principal component')
plt.title('bar plot')
plt.show()
# 7 component is enough for model.
```



In [36]:

```
lm=LinearRegression()
pcr_model=lm.fit(X_reduced_train,y_train)
print('Intercept: ' + str(pcr_model.intercept_))
print('Coefficients: ' + str(pcr_model.coef_))
```

```
Intercept: -0.005461278328146827
Coefficients: [-0.13248389  0.18798958  0.64978531 -0.24235563 -0.038966   0.34814071
 -0.01512341 -0.38662935 -0.30625748  0.14768885  0.59449755 -3.03336011]
```

In [37]:

```
#Prediction
y_pred=pcr_model.predict(X_reduced_train)
np.sqrt(mean_squared_error(y_train,y_pred))
```

Out[37]:

```
0.268553922694838
```

In [38]:

```
df['motor_speed'].mean()
```

Out[38]:

```
-0.006335507987812318
```

In [39]:

```
#R squared
r2_score(y_train,y_pred)
```

Out[39]:

```
0.9280615197697045
```

In [40]:

```
# Prediction For testing error
pca2=PCA()

X_reduced_test=pca2.fit_transform(scale(X_test))
```

```

pcr_model2=lm.fit(X_test,y_test)

y_pred=pcr_model2.predict(X_reduced_test)

print('RMSE for test model : ' +str(np.sqrt(mean_squared_error(y_test,y_pred))))

```

RMSE for test model : 1.657960891802902

In [41]:

```

#Model Tuning for PCR

lm=LinearRegression()
pcr_model=lm.fit(X_reduced_train[:,0:10],y_train)
y_pred=pcr_model.predict(X_reduced_test[:,0:10])

from sklearn import model_selection

cv_10=model_selection.KFold(n_splits=10,
                             shuffle=True,
                             random_state=1)

```

In [44]:

```

lm=LinearRegression()
RMSE=[]

for i in np.arange(1,X_reduced_train.shape[1] + 1):
    score=np.sqrt(-1*model_selection.cross_val_score(lm,
                                                       X_reduced_train[:,i],
                                                       y_train.ravel(),
                                                       cv=cv_10,
                                                       scoring='neg_mean_squared_error').mean())

    RMSE.append(score)


```

In [43]:

```

plt.plot(RMSE)
plt.xlabel('# of Components')
plt.ylabel('RMSE')
plt.title('PCR Model Tuning for Motor_Speed Prediction');

```



In [45]:

```

##10 component is good for the model because RMSE value is the smallest for this component number.
##That's why there is no need to tune the model.

```

In [6]:

```

*****Polynomial Regression*****
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import PolynomialFeatures

```

```

from sklearn.preprocessing import PolynomialFeatures
from sklearn.metrics import r2_score, mean_squared_error
import pandas as pd

```

In [7]:

```

df=pd.read_csv('C:\\Users\\Mahima Sharu\\pmsm_temperature_data.csv')
X=df.drop("motor_speed", axis=1)
y=df["motor_speed"]
from sklearn.model_selection import train_test_split, cross_val_score, cross_val_predict
X_train, X_test, y_train, y_test=train_test_split(X, y, test_size=0.2, random_state=42)
training=df.copy()

quad = PolynomialFeatures (degree = 2)
x_quad = quad.fit_transform(X_train)

X_train, X_test, y_train, y_test = train_test_split(x_quad, y_train, random_state = 0)

plr = LinearRegression().fit(X_train, y_train)

Y_train_pred = plr.predict(X_train)
Y_test_pred = plr.predict(X_test)

print('Polynomial Linear Regression: ', plr.score(X_test, y_test))

```

Polynomial Linear Regression: 0.9952573854207651

In [10]:

```

#Plotting Residual in Linear Regression

import matplotlib.pyplot as plt
from sklearn import linear_model, metrics
#Create linear regression object
reg=linear_model.LinearRegression()

#train the model using the train data sets
reg.fit(X_train, y_train)

#regression coefficients
print("Coefficients: \n", reg.coef_)

#Variance score
print("Variance score: {}".format(reg.score(X_test, y_test)))

plt.style.use('fivethirtyeight')

#plotting residual errors in training data
plt.scatter(reg.predict(X_train), reg.predict(X_train)-y_train,
            color="green", s=10, label="train data")

#plotting residual errors in test data
plt.scatter(reg.predict(X_test), reg.predict(X_test)-y_test,
            color="blue", s=10, label="test data")

#plot line for zero residual error
plt.hlines(y=0, xmin=-2, xmax=2, linewidth=2)

#plot legend
plt.legend(loc='upper right')

#plot title
plt.title("residual error")

plt.show()

```

Coefficients:


```

[-1.19652298e-12 -1.94751304e-03  8.86767668e-02 -1.51002720e-01
 8.46456963e-01 -4.33513690e-01 -4.11436417e-01  2.67644301e-01
 1.87790064e-02 -1.25988463e-01 -1.71853359e-01  2.54897914e-01
 7.40823232e-04 -4.34110557e-03 -3.24485408e-02  1.34249919e-02
-1.05690117e-04  1.67560244e-01  8.52062538e-03 -1.56724820e-01
-1.89292501e-03  5.46155535e-02 -9.53354640e-03 -1.36632243e-02
 2.28461480e-05  6.35750499e-02 -1.63781265e-01 -5.32017845e-02

```

```
-9.07557261e-01 -5.67804823e-02 7.42197288e-01 1.03560119e-01
-4.73717311e-01 3.83839811e-01 -4.06039694e-02 -6.25699752e-04
1.38771773e-01 3.68937154e-02 -8.83271855e-01 2.50808967e-02
1.04897761e+00 -3.74045176e-02 6.23406019e-01 -6.54857489e-01
1.79516489e-01 -1.55441855e-03 -4.59404937e-05 -3.55526844e-01
-3.64804458e-01 2.34531508e-01 -3.94524138e-03 2.36891662e-01
-2.88885594e-01 1.08053335e-01 -3.56503666e-04 7.56339301e-01
7.49643580e-01 -3.33692091e+00 -3.06318572e-01 4.23305718e+00
-5.39802183e+00 2.10899507e+00 -3.84090475e-03 3.82247855e-01
-6.62559679e-01 -3.21686761e-02 2.04955803e-01 -2.70200980e-01
1.16040888e-01 -3.21988930e-04 2.61534726e+00 2.63818883e-01
-3.58802305e+00 4.68164431e+00 -1.89010204e+00 2.42748089e-03
8.90538096e-03 -3.41181602e-01 3.53620032e-01 -1.06780573e-01
2.51351405e-04 8.80366838e-01 -1.46367360e+00 2.09986581e-01
6.14259740e-04 5.86532523e-01 -1.59061514e-01 3.16715435e-03
9.53273810e-03 -3.66373716e-03 -5.86982787e-06]
```

Variance score: 0.9952573854207651



In []: